

For a DevOps-based role specializing in AWS Lambda, you'll want to focus on the following key topics:

Key AWS Lambda Topics to Focus On:

1. Lambda Basics:

- What is AWS Lambda?
- Event-driven architecture
- AWS Lambda pricing (Pay-as-you-go model)
- Cold starts and warm invocations

Cold starts = slower, due to setup and initialization, Warm invocations = faster, reusing existing environment. To mitigate cold start AWS gives "provisioned concurrency," which keeps a set number of execution environments pre-warmed and Minimize heavy initialization logic, such as opening database connections or loading large files, to reduce the impact of cold starts.

- Limits and quotas (memory, execution timeout, etc.)

Maximum Timeout: 15 minutes (900 seconds)

Default: 3 seconds

Default Limit: 1,000 concurrent executions (can be increased with AWS support)

Provisioned Concurrency: Ensures warm starts for a set number of concurrent instances.

Ephemeral Disk Storage: 512 MB (temporary file storage during execution)

Max File Descriptors/Processes: 1,024 file descriptors, 1,024 processes/threads

2. Integration with Other AWS Services:

- Lambda triggers (S3, SNS, SQS, DynamoDB, CloudWatch, etc.)
- Lambda with API Gateway
- Lambda with Step Functions for orchestration
- Lambda with S3 events, DynamoDB streams, etc.

3. Deployment & Automation:

- AWS SAM (Serverless Application Model) or CloudFormation for deploying Lambda functions
- Infrastructure as Code (Terraform/CloudFormation) for automating Lambda deployments
- CI/CD for Lambda with tools like Jenkins, CodePipeline

4. Lambda Layers:

Lambda

- Reusable libraries and dependencies
- Packaging and using layers

5. Monitoring & Logging:

- Lambda integration with CloudWatch for logs and monitoring
- X-Ray for tracing
- Alarms and notifications

6. Security:

- IAM roles and policies for Lambda execution
- VPC integration for Lambda functions
- Securing Lambda with KMS (Key Management Service)

7. Advanced Concepts:

- Concurrency and throttling
- Error handling (retry mechanisms, dead-letter queues)
- Provisioned concurrency for reducing cold starts
- Lambda destinations (asynchronous invocations)
- Performance optimization techniques

Basic Template for Writing an AWS Lambda Function (Python):

```
import json

def lambda_handler(event, context):
    # Process the event data
    response = {
        'statusCode': 200,
        'body': json.dumps('Hello from Lambda!')
    }
    return response
```

Template Breakdown:

- `lambda_handler(event, context)`: The entry point for your function.
- `event`: The data passed to your function (triggered by another AWS service).
- `context`: Provides runtime information like request IDs, timeout, etc.
- `Response`: The return value from the Lambda function that includes the HTTP status and body.

AWS Lambda Real-Time Project-Based Examples:

1. Cost Optimization:

Goal: Automatically stop idle EC2 instances to save costs.

```
import boto3
def lambda_handler(event, context):
    ec2 = boto3.client('ec2')
    instances = ec2.describe_instances(Filters=[{'Name': 'instance-state-name', 'Values': ['running']}])

    for reservation in instances['Reservations']:
        for instance in reservation['Instances']:
            instance_id = instance['InstanceId']
            ec2.stop_instances(InstanceIds=[instance_id])
            print(f'Stopped EC2 instance: {instance_id}')
```

2. Billing Management:

Goal: Send an alert if the AWS monthly cost exceeds a certain threshold.

```
import boto3
import json
from datetime import datetime

def lambda_handler(event, context):
    cloudwatch = boto3.client('cloudwatch')
    billing_alarm = cloudwatch.describe_alarms(AlarmNames=['BillingAlarm'])

    if billing_alarm['MetricAlarms'][0]['StateValue'] == 'ALARM':
        print("Billing threshold exceeded. Notify stakeholders.")
        # You can integrate with SNS for notifications here.
```

3. Resource Cleanup:

Goal: Automatically delete old EBS snapshots that are older than 30 days.

```
import boto3
from datetime import datetime, timedelta

def lambda_handler(event, context):
    ec2 = boto3.client('ec2')
    retention_days = 30
    delete_time = datetime.now() - timedelta(days=retention_days)

    snapshots = ec2.describe_snapshots(OwnerIds=['self'])['Snapshots']
    for snapshot in snapshots:
        snapshot_time = snapshot['StartTime'].replace(tzinfo=None)
```

```
if snapshot_time < delete_time:
    ec2.delete_snapshot(SnapshotId=snapshot['SnapshotId'])
    print(f'Deleted snapshot: {snapshot["SnapshotId"]})
```

4. Automated Backup:

Goal: Create daily snapshots for all running EC2 instances.

```
import boto3
def lambda_handler(event, context):
    ec2 = boto3.client('ec2')
    instances = ec2.describe_instances(Filters=[{'Name': 'instance-state-name', 'Values': ['running']}])

    for reservation in instances['Reservations']:
        for instance in reservation['Instances']:
            instance_id = instance['InstanceId']
            ec2.create_snapshot(Description=f'Snapshot of {instance_id}', InstanceId=instance_id)
            print(f'Created snapshot for EC2 instance: {instance_id})
```

5. Log Processing:

Goal: Automatically process and store logs from an S3 bucket to a DynamoDB table.

```
import boto3
def lambda_handler(event, context):
    s3 = boto3.client('s3')
    dynamodb = boto3.resource('dynamodb')
    table = dynamodb.Table('LogTable')

    for record in event['Records']:
        bucket = record['s3']['bucket']['name']
        key = record['s3']['object']['key']

        log_data = s3.get_object(Bucket=bucket, Key=key)
        log_content = log_data['Body'].read().decode('utf-8')

        table.put_item(Item={'log_key': key, 'log_data': log_content})
        print(f'Stored log {key} into DynamoDB')
```

6. Monitoring and Alerts:

Goal: Create a Lambda function to monitor CPU usage of EC2 instances and trigger alerts.

```
import boto3
def lambda_handler(event, context):
    cloudwatch = boto3.client('cloudwatch')
    sns = boto3.client('sns')

    alarms = cloudwatch.describe_alarms(
```

```
        AlarmNames=['HighCPUUsage'],  
        StateValue='ALARM'  
    )  
  
    if alarms['MetricAlarms']:  
        sns.publish(TopicArn='arn:aws:sns:region:account-id:topic', Message='High CPU Usage Alert!')  
        print("High CPU alert sent!")
```

Categories for Quick Revision:

1. Cost Optimization:

- EC2 Idle Instance Management
- Unused Resources Cleanup (e.g., EBS snapshots)

2. Billing Management:

- Billing Alerts
- Cost Threshold Monitoring

3. Resource Automation:

- Automated EC2 Backups
- Scheduled Scaling & Resource Management

4. Log Processing:

- S3 log data processing
- Storing logs in DynamoDB

5. Monitoring & Alerts:

- CPU Usage Monitoring
- Event-driven alerts

6. Security Automation:

- IAM role enforcement
- VPC network security

By focusing on these topics and examples, you'll be well-prepared for your AWS Lambda-based DevOps interview!